

10-701: Introduction to

~~Deep Neural Networks~~

Machine Learning

<http://www.cs.cmu.edu/~10701>


Organizational info

- All up-to-date info is on the course web page (follow links from my page).
- Instructors
 - Nina balcan
 - Ziv Bar-Joseph
- TAs: See info on website for recitations, office hours etc.
- See web page for contact info, office hours, etc.
- Piazza would be used for questions / comments and likely for class quizzes. Make sure you are subscribed.

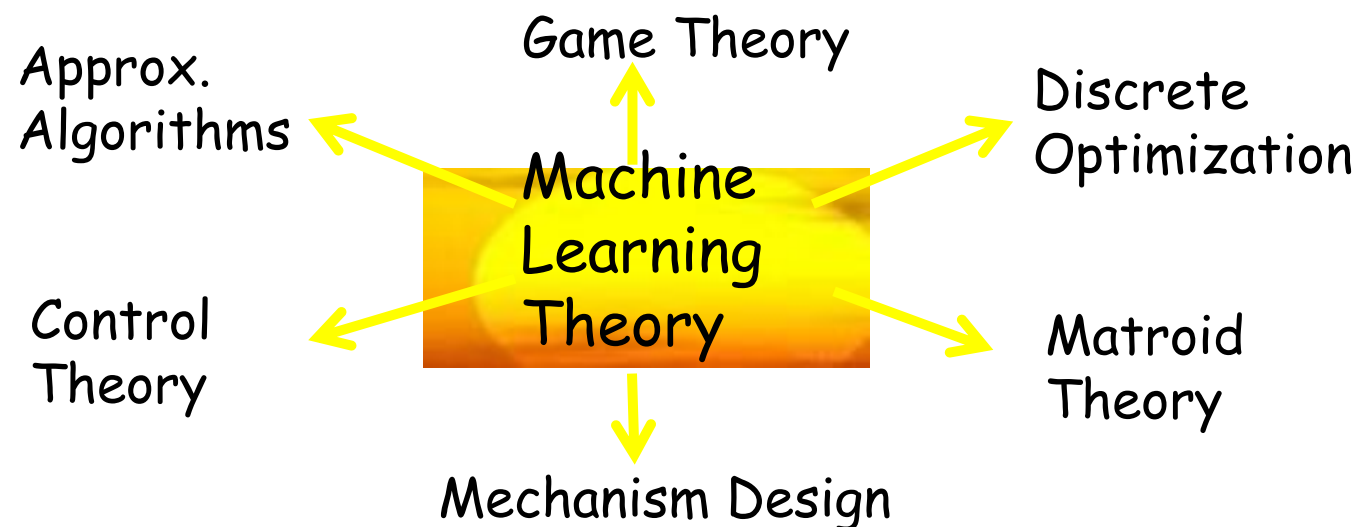
Maria-Florina Balcan: Nina



- Foundations for Modern Machine Learning
- E.g., interactive, semi-supervised, distributed, multi-task, never-ending, privacy preserving learning



- Connections between learning theory & other fields (algorithms, algorithmic game theory)



- Program Committee Chair for ICML 2016 (main general machine learning conference), COLT 2014 (main learning theory conference)

Sarah Schultz (Assistant Lecturer)

sschultz@cs.cmu.edu

GHC 8110

Research Interests:

Educational data mining and
Intelligent Tutoring Systems





Ellen Vitercik

Email: vitercik@cs.cmu.edu

Office hours: Friday 10-11 in GHC 7511

Research interests:


Theoretical machine learning

Computational economics

Logan Brooks (lcbrooks@andrew)



- Office space: GHC 6219
- Office hours: Monday 10-11
- Research topic: epidemic forecasting
 - Time series
 - Ensembles



Yujie Xu (yujie.xu@andrew.cmu.edu)



- GHC 5th floor common area near entrance
- Office Hours: Mon 4:30-5:30
- Research topic:
data-driven building energy models
 - regression
 - impact evaluation

Easwaran Ramamurthy

eramamur@andrew.cmu.edu



- Find Me: GHC 7405
- Office Hours: Tuesday 4-5
- Interests:
 - Computational Genomics
 - Deep learning applications in regulatory genomics
 - Alzheimer's Disease

Chieh Lin

(chiehl1@cs.cmu.edu)

Office:


GHC 8021

Office Hours:

Thursday 10:30-11:30

Research Interest:

1. ML Applications in biological/medical Data
2. Neural Networks/Deep Learning



Matt Oresky

(moresky@andrew.cmu.edu)



Office Hours:
Tuesday 9:30 – 10:30 AM
GHC 6th floor common area
(by the kitchenette)

Interest:
Natural Language Processing

Akash Ramachandran

(akashr1@andrew.cmu.edu)



- Office hours : Friday 3-4pm
- Interests:
 - Application of ML in Biology
 - Software Development in Java
 - Playing the *tabla* (an Indian drum)

Guoquan Zhao

(guoquanz@andrew.cmu.edu)



Find me: GHC 6th floor common area

Office Hours:

Thursday 3.30 pm – 4.30 pm

Interest:

Active Learning

Distributed ML system

- 8/28 Introduction, MLE
- 8/30 Classification, KNN
- 9/4 – no class, labor day
- 9/6 – Decision trees / problem set 1 out
- 9/11 – Naïve Bayes
- 9/13 – Linear regression
- 9/18 – Logistic regression
- 9/20 – Graphical Models, MRF/ PS1 due, PS2 out
- 9/25 – Gr **10/25 (Wednesday): Midterm**
- 9/27 – Graphical Models, BN 2
- 10/2 – Perceptron
- 10/4 – Kernel Methods/ PS2 due, PS3 out
- 10/9 – Support Vector Machines
- 10/11 – Neural networks 1: Backpropagation
- 10/16– Neural networks 2: Deep NN/ project proposals
- 10/18 – Ensemble Learning, Boosting / PS3 due
- 10/23 – Active Learning
- 10/25 Midterm/ PS4 out
- 10/30 – Dimensionality Reduction
- 11/1 – Unsupervised learning (clustering)
- 11/6 – Semi supervised learning
- 11/8 - Generalization, overfitting I / PS 4 due, PS 5 out
- 11/13 – Model Selection.
- 11/15 – Hidden markov models – learning
- 11/20 – HMM – inference
- 11/22 – no class, thanksgiving break
- 11/27 – MDPS
- 11/29 –Reinforcement Learning / PS 5 due
- 12/4 – Distributed ML?
- 12/6 – Final review

Intro and classification (A.K.A. 'supervised learning')

Graphical models

Non linear and kernel methods

Unsupervised learning

• Theoretical considerations

Reasoning under uncertainty

Grading

- **5 Problem sets (5th has a higher weight) - 45%**
- **Final - 30%**
- **Midterm - 20%**
- **Class participation - 5%**

Class assignments

- 5 Problem sets
 - Most containing both theoretical and programming assignments
 - Last problems set: mini project
- Exams
 - Midterm (10/25)
 - Final

Recitations

- Twice a week (same content in both)
- Expand on material learned in class, go over problems from previous classes etc.

What is Machine Learning?

Easy part: Machine

Hard part: Learning

- Short answer: Methods that can help generalize information from the observed data so that it can be used to make better decisions in the future

What is Machine Learning?

Longer answer: The term Machine Learning is used to characterize a number of different approaches for generalizing from observed data:

- Supervised learning
 - Given a set of features and labels learn a model that will predict a label to a new feature set
- Unsupervised learning
 - Discover patterns in data
- Reasoning under uncertainty
 - Determine a model of the world either from samples or as you go along
- Active learning
 - Select not only model but also which examples to use

Paradigms of ML

- Supervised learning
 - Given $D = \{X_i, Y_i\}$ learn a model (or function) $F: X_k \rightarrow Y_k$
- Unsupervised learning
 - Given $D = \{X_i\}$ group the data into Y classes using a model (or function) $F: X_i \rightarrow Y_j$
- Reinforcement learning (reasoning under uncertainty)
 - Given $D = \{\text{environment, actions, rewards}\}$ learn a policy and utility functions:

policy: $F1: \{e, r\} \rightarrow a$
utility: $F2: \{a, e\} \rightarrow R$
- Active learning
 - Given $D = \{X_i, Y_i\}, \{X_j\}$ learn a function $F1: \{X_j\} \rightarrow x_k$ to maximize the success of the supervised learning function $F2: \{X_i, x_k\} \rightarrow Y$

Recommender systems

Amazon.com: Recommended for You - Mozilla Firefox

File Edit View Go Bookmarks Tools Help

http://www.amazon.com/gp/yourstore/ref=pd_irl_283155?ie=UTF8&nodeID=283155&rGroup=books&pf_rd_p=273115801&pf_rd_s=center-2&pf_rd_t=101&pf_rd_i=2831558

Getting Started Latest Headlines

amazon.com Hello, Ziv Bar-Joseph. We have recommendations for you. (Not Ziv?)

Get FREE Two-Day Shipping Now

Ziv's Amazon.com Today's Deals Gifts & Wish Lists Gift Cards Your Account | Help

Shop All Departments Search Books GO Cart Your Lists

Ziv's Amazon.com Your Browsing History Recommended For You Rate These Items Improve Your Recommendations Your

Recommended For You > Books

Recommendations by Category in Books

Featured Categories Refinements Special Features Specialty Stores Books

Improve Your Recommendations


Update your Amazon history to improve your recommendations

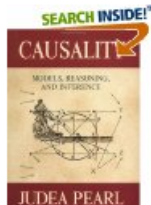
Items you own (13) Rated items Not Interested

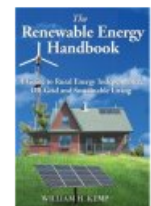
Need Help? Visit our help area to learn more.


These recommendations are based on [items you own](#) and more.

view: All | New Releases | Coming Soon

1.  **Pattern Recognition and Machine Learning (Information Science and Statistics)**
by Christopher M. Bishop (Oct 1, 2007)
Average Customer Review: ★★★★★ (38)
In Stock
List Price: \$84.95
Price: \$62.60
56 used & new from \$56.64
☐ I own it ☐ Not interested ☒ ★★★★★ Rate it
Recommended because you purchased **Learning in Graphical Models** and more ([Fix this](#))

2.  **Causality: Models, Reasoning, and Inference**
by Judea Pearl (Mar 13, 2000)
Average Customer Review: ★★★★★ (12)
In Stock
List Price: \$50.00
Price: \$38.50
26 used & new from \$32.01
☐ I own it ☐ Not interested ☒ ★★★★★ Rate it
Recommended because you purchased **Probabilistic Reasoning in Intelligent Systems** and more ([Fix this](#))

3.  **The Renewable Energy Handbook: A Guide to Rural Energy Independence, Off-Grid and Sustainable Living**
by William H. Kemp (April 1, 2006)
Average Customer Review: ★★★★★ (16)
In Stock
List Price: \$29.95
Price: \$19.77
40 used & new from \$18.25
☐ I own it ☐ Not interested ☒ ★★★★★ Rate it
Recommended because you purchased **Wind Power, Revised Edition** and more ([Fix this](#))

4.  **Learning Bayesian Networks (Artificial Intelligence)**
by Richard E. Neapolitan (April 6, 2003)
Average Customer Review: ★★★★★ (2)

http://www.amazon.com/Pattern-Recognition-Learning-Information-Statistics/dp/0387310738/ref=pd_ys_ir_b_1?pf_rd_p=258372101&pf_rd_s=center-1&pf_rd_t=1501&pf_rd_i=list&pf_rd_m=ATVPDKIKX0DER&pf_rd_r=1BQMM558P495ESDQ9BHP

start Inbox for zivbj@cs... C:\ziv\classes\AI08... 3 Microsoft Power... preInqAuthor - Mic... Carnegie Mellon Uni... Amazon.com: Reco... 4:28 PM

Primarily supervised learning

NELL: Never-Ending Language Learning

Can computers learn to read? We think so. "Read the Web" is a research project that attempts to create a computer system that learns over time to read the web. Since January 2010, our computer system called NELL (Never-Ending Language Learner) has been running continuously, attempting to perform two tasks each day:

- First, it attempts to "read," or extract facts from text found in hundreds of millions of web pages (e.g., `playsInstrument(George_Harrison, guitar)`).
- Second, it attempts to improve its reading competence, so that tomorrow it can extract more facts from the web, more accurately.

So far, NELL has accumulated over 50 million candidate beliefs by reading the web, and it is considering these with confidence. NELL has high confidence in 3,938,530 of these beliefs — these are displayed on this website. It is not perfect, but NELL is learning. You can track NELL's progress below or [@cmunell on Twitter](#), browse and download its [knowledge base](#), read more about our [technical approach](#), or join the [discussion group](#).



semi supervised
learning

Recently-Learned Facts



Refresh

instance	iteration	date learned	confidence	
glass_window_restoration is a household item	1069	03-aug-2017	97.5	
bracelets_curb is a kind of clothing	1069	03-aug-2017	90.9	
hillsborough_lista_d_attesa_crea_un_gruppo_meetup is a visualizable thing	1069	03-aug-2017	99.1	
parison_levitra_viagra_cialis is a drug	1069	03-aug-2017	97.7	
the_democratic_daily is a newspaper	1069	03-aug-2017	100.0	
barcelona_international_airport is an airport in the city barcelona	1073	22-aug-2017	100.0	
john003 has brother james	1073	22-aug-2017	100.0	
omaha_world_herald is a newspaper in the city new_york	1073	22-aug-2017	93.8	
abc is a company headquartered in the city new_york	1073	22-aug-2017	100.0	
arachnids001 is an arthropod as well as mites also is	1073	22-aug-2017	93.8	

Driveless cars

Supervised and
reinforcement learning

Helicopter control



Reinforcement learning

Deep neural networks

Supervised learning (though
can also be trained in an
unsupervised way)

Distributed gradient descent based on bacterial movement

Reasoning under
uncertainty



Biology

ACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAACGCTGAGCAATTCGATAGCAATTC
GATAACGCTGAGCAATTCGGATAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAACG
CTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGGATATCGATAGCAATTCGATAAATC
GGATAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAACGCTGAGCAATTCGATAGC
AATTCGATAACGCTGAGCAATTCGGATATCGATAGCAATTCGATAACGCTGAGCAACGCTGAGCA
ATTCGATAGCAATTCGATAACGCTGAGCAATTCGGATAACGCTGAGCAATTCGATAGCATTCGAT
AACGCTGAGCAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGGATAACGCTG
AGCAATTCGATAGCAATTCGATAACGCTGAGCTGAGCAATTCGATAGCAATTCGATAACGCTGA
GCAATTCGGATAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGATAACGCTG
GATAGCAATTCGATAGCAATTCGATAGCAATTCGATAACGCTGAGCAACGCTGAGCAATTCGAT
AGCAATTCGATAACGCTGAGCAATTCGGATAACGCTGAGCAATTCGATAGCAATTCGATAACGCT
GAGCAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGGATATCGATAGCAATT
CGATAACGCTGAGCAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGGATAAC
GCTGAGCAATTCGATAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGATAACGCTGAG
CTGATAGCAATTCGATAACGCTGAGCAACGCTGAGCAACGCTGAGCA
ATTCGGATATCGATAGCAATTCGATAACGCTGAGCA
ACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGGATAACGCTGAGCAATTCGAT
AGCATTTCGATAACGCTGAGCAACGCTGAGCAATTCGATAGCAATTCGATCGGATAACGCTGAGC
AATTCGATAGCAATTCGATAACGCTGAGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCA
ATCGGATAACGCTGAGCAATTCGATAGCA
AGCAATTCGATAACGCTGAGCAATTCGGAT
GCAATTCGATAGCAATTCGATAACGCTGA
GATAACGCTGAGCAACGCTGAGCAATTCG
CTGAGCAATTCGATAGCAATTCGATAACG
TGAGCAATTCGGATAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAACGCTGAGCAA
TTCGATAGCAATTCGATAGCAATTCGATAGCAATTCGATAACGCTGAGCAACGCTGAGCAATTC
GATAGCAATTCGATAACGCTGAGCAATTCGGATAACGCTGAGCAATTCGATAGCAATTCGATAAC
GCTGAGCAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGGATATCGATAGCA
ATTCGATAACGCTGAGCAACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCAATTCGGAT
AACGCTGAGCAATTCGATAGCAATTCGATAACGCTGAGCTGAGCAATTCGATAGCAATTCGATA
ACGCTGAGCAATTCGGA

Which part is the gene?

Supervised and
unsupervised learning (can
also use active learning)

Common Themes

- Mathematical framework
 - Well defined concepts based on explicit assumptions
- Representation
 - How do we encode text? Images?
- Model selection
 - Which model should we use? How complex should it be?
- Use of prior knowledge
 - How do we encode our beliefs? How much can we assume?

(brief) intro to probability

Basic notations

- Random variable
 - referring to an element / event whose status is unknown:
 $A = \text{"it will rain tomorrow"}$
- Domain (usually denoted by Ω)
 - The set of values a random variable can take:
 - " $A = \text{The stock market will go up this year}$ ": Binary
 - " $A = \text{Number of Steelers wins in 2015}$ ": Discrete
 - " $A = \text{\% change in Google stock in 2015}$ ": Continuous

Axioms of probability (Kolmogorov's axioms)

A variety of useful facts can be derived from just three axioms:

1. $0 \leq P(A) \leq 1$
2. $P(\text{true}) = 1$, $P(\text{false}) = 0$
3. $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

There have been several other attempts to provide a foundation for probability theory. Kolmogorov's axioms are the most widely used.

Priors

Degree of belief
in an event in the
absence of any
other information

No rain




$$P(\text{rain tomorrow}) = 0.2$$

$$P(\text{no rain tomorrow}) = 0.8$$


Conditional probability

- $P(A = 1 \mid B = 1)$: The fraction of cases where A is true if B is true

$$P(A = 0.2)$$



$$P(A|B = 0.5)$$



Conditional probability

- In some cases, given knowledge of one or more random variables we can improve upon our prior belief of another random variable
- For example:

$$p(\text{slept in movie}) = 0.5$$

$$p(\text{slept in movie} \mid \text{liked movie}) = 1/4$$

$$p(\text{didn't sleep in movie} \mid \text{liked movie}) = 3/4$$

Slept	Liked
1	0
0	1
1	1
1	0
0	0
1	0
0	1
0	1

Joint distributions

- The probability that a *set* of random variables will take a specific value is their joint distribution.
- Notation: $P(A \wedge B)$ or $P(A,B)$
- Example: $P(\text{liked movie, slept})$

If we assume independence then

$$P(A,B)=P(A)P(B)$$

However, in many cases such an assumption may be too strong
(more later in the class)

Joint distribution (cont)

$P(\text{class size} > 20) = 0.6$

$P(\text{summer}) = 0.4$

$P(\text{class size} > 20, \text{summer}) = ?$

Evaluation of classes

Size	Time	Eval
30	R	2
70	R	1
12	S	2
8	S	3
56	R	1
24	S	2
10	S	3
23	R	3
9	R	2
45	R	1

Joint distribution (cont)

$P(\text{class size} > 20) = 0.6$

$P(\text{summer}) = 0.4$

$P(\text{class size} > 20, \text{summer}) = 0.1$

Evaluation of classes

Size	Time	Eval
30	R	2
70	R	1
12	S	2
8	S	3
56	R	1
24	S	2
10	S	3
23	R	3
9	R	2
45	R	1

Joint distribution (cont)

$P(\text{class size} > 20) = 0.6$

$P(\text{eval} = 1) = 0.3$

$P(\text{class size} > 20, \text{eval} = 1) = 0.3$

Size	Time	Eval
30	R	2
70	R	1
12	S	2
8	S	3
56	R	1
24	S	2
10	S	3
23	R	3
9	R	2
45	R	1

Joint distribution (cont)

$P(\text{class size} > 20) = 0.6$

$P(\text{eval} = 1) = 0.3$

$P(\text{class size} > 20, \text{eval} = 1) = 0.3$

Evaluation of classes


Size	Time	Eval
30	R	2
70	R	1
12	S	2
8	S	3
56	R	1
24	S	2
10	S	3
23	R	3
9	R	2
45	R	1

Chain rule

- The joint distribution can be specified in terms of conditional probability:

$$P(A,B) = P(A|B) \cdot P(B)$$

- Together with Bayes rule (which is actually derived from it) this is one of the most powerful rules in probabilistic reasoning



Bayes rule

- One of the most important rules for this class.
- Derived from the chain rule:

$$P(A,B) = P(A | B)P(B) = P(B | A)P(A)$$

- Thus,

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$




Thomas Bayes was an English clergyman who set out his theory of probability in 1764.


Bayes rule (cont)

Often it would be useful to derive the rule a bit further:


$$P(A|B) = \frac{P(B|A)P(A)}{P(B)} = \frac{P(B|A)P(A)}{\sum_A P(B|A)P(A)}$$



$P(B, A=1)$



$P(B, A=0)$

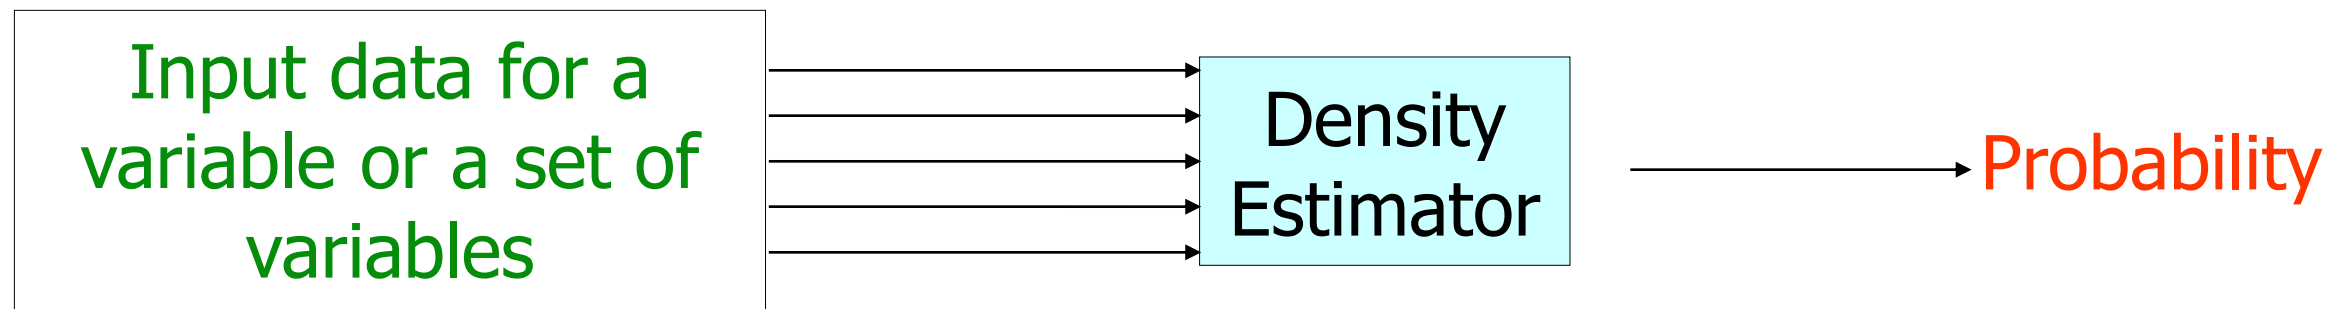


This results from:
 $P(B) = \sum_A P(B, A)$

Density estimation

Density Estimation

- A Density Estimator learns a mapping from a set of attributes to a Probability



Density estimation

- Estimate the distribution (or conditional distribution) of a random variable
- Types of variables:
 - Binary
coin flip, alarm
 - Discrete
dice, car model year
 - Continuous
height, weight, temp.,

When do we need to estimate densities?

- Density estimators are critical ingredients in several of the ML algorithms we will discuss
- In some cases these are combined with other inference types for more involved algorithms (i.e. EM) while in others they are part of a more general process (learning in BNs and HMMs)

Density estimation

- Binary and discrete variables:

Easy: Just count!

- Continuous variables:

Harder (but just a bit): Fit
a model

Learning a density estimator for discrete variables

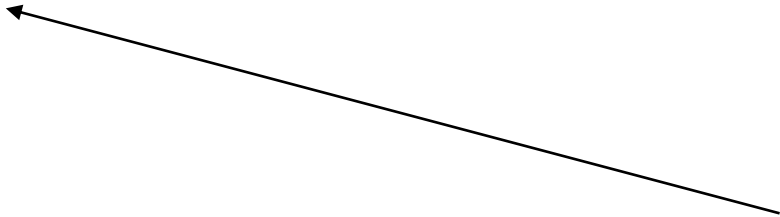
$$\hat{P}(x_i = u) = \frac{\text{\#records in which } x_i = u}{\text{total number of records}}$$

A trivial learning algorithm!

But why is this true?

Maximum Likelihood Principle

We can define the likelihood of the data given the model as follows:

$$\hat{P}(\text{dataset} \mid M) = \hat{P}(x_1 \wedge x_2 \dots \wedge x_n \mid M) = \prod_{k=1}^n \hat{P}(x_k \mid M)$$


M is our model (usually a collection of parameters)

For example M is

- The probability of 'head' for a coin flip
- The probabilities of observing 1,2,3,4 and 5 for a dice
- etc.

Maximum Likelihood Principle

$$\hat{P}(\text{dataset} \mid M) = \hat{P}(x_1 \wedge x_2 \dots \wedge x_n \mid M) = \prod_{k=1}^n \hat{P}(x_k \mid M)$$

- Our goal is to determine the values for the parameters in M
- We can do this by maximizing the probability of generating the observed samples
- For example, let Θ be the probabilities for a coin flip
- Then

$$L(x_1, \dots, x_n \mid \Theta) = p(x_1 \mid \Theta) \dots p(x_n \mid \Theta)$$

- The observations (different flips) are assumed to be independent
- For such a coin flip with $P(H)=q$ the best assignment for Θ_h is

$$\operatorname{argmax}_q = \#H/\#\text{samples}$$

- Why?

Maximum Likelihood Principle: Binary variables


- For a binary random variable A with $P(A=1)=q$
 $\operatorname{argmax}_q = \#1/\#\text{samples}$

- Why?

Data likelihood: $P(D | M) = q^{n_1} (1 - q)^{n_2}$

We would like to find: $\operatorname{argmax}_q q^{n_1} (1 - q)^{n_2}$

Omitting terms that
do not depend on q



Maximum Likelihood Principle

Data likelihood: $P(D | M) = q^{n_1} (1 - q)^{n_2}$

We would like to find: $\arg \max_q q^{n_1} (1 - q)^{n_2}$

$$\frac{\partial}{\partial q} q^{n_1} (1 - q)^{n_2} = n_1 q^{n_1-1} (1 - q)^{n_2} - q^{n_1} n_2 (1 - q)^{n_2-1}$$

$$\frac{\partial}{\partial q} = 0 \Rightarrow$$

$$n_1 q^{n_1-1} (1 - q)^{n_2} - q^{n_1} n_2 (1 - q)^{n_2-1} = 0 \Rightarrow$$

$$q^{n_1-1} (1 - q)^{n_2-1} (n_1 (1 - q) - q n_2) = 0 \Rightarrow$$

$$n_1 (1 - q) - q n_2 = 0 \Rightarrow$$

$$n_1 = n_1 q + n_2 q \Rightarrow$$

$$q = \frac{n_1}{n_1 + n_2}$$


Log Probabilities

When working with products, probabilities of entire datasets often get too small. A possible solution is to use the log of probabilities, often termed 'log likelihood'

$$\log \hat{P}(\text{dataset} \mid M) = \log \prod_{k=1}^n \hat{P}(x_k \mid M) = \sum_{k=1}^n \log \hat{P}(x_k \mid M)$$

Maximizing this likelihood function is the same as maximizing $P(\text{dataset} \mid M)$

Log values
between 0 and 1



In some cases moving to log space would also make computation easier (for example, removing the exponents)

How much do grad students sleep?

- Lets try to estimate the distribution of the time students spend sleeping (outside class).

Possible statistics

- **X**

Sleep time

- **Mean of X :**


$$E\{X\}$$

7.03

- **Variance of X :**

$$\text{Var}\{X\} = E\{(X - E\{X\})^2\}$$


3.05



Covariance: Sleep vs. GPA

- Co-Variance of X1,
X2:

$$\begin{aligned} \text{Covariance}\{X1, X2\} &= \\ E\{(X1 - E\{X1\})(X2 - E\{X2\})\} &= \\ &= 0.88 \end{aligned}$$



Statistical Models

- Statistical models attempt to characterize properties of the population of interest
- For example, we might believe that repeated measurements follow a normal (Gaussian) distribution with some mean μ and variance σ^2 , $x \sim N(\mu, \sigma^2)$


where

$$p(x | \Theta) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(x-\mu)^2}{2\sigma^2}}$$

and $\Theta=(\mu, \sigma^2)$ defines the parameters (mean and variance) of the model.


The Parameters of Our Model

- A statistical model is a **collection** of distributions; the **parameters** specify individual distributions $x \sim N(\mu, \sigma^2)$
- We need to adjust the parameters so that the resulting distribution **fits** the data well




The Parameters of Our Model

- A statistical model is a **collection** of distributions; the **parameters** specify individual distributions $x \sim N(\mu, \sigma^2)$
- We need to adjust the parameters so that the resulting distribution **fits** the data well



Computing the parameters of our model

- Lets assume a Gaussian distribution for our sleep data
- How do we compute the parameters of the model?



Maximum Likelihood Principle

- We can fit statistical models by maximizing the probability of generating the observed samples:

$$L(x_1, \dots, x_n \mid \Theta) = p(x_1 \mid \Theta) \dots p(x_n \mid \Theta)$$

(the samples are assumed to be independent)

- In the Gaussian case we simply set the mean and the variance to the sample mean and the sample variance:

$$\overline{\mu} = \frac{1}{n} \sum_{i=1}^n x_i \qquad \overline{\sigma^2} = \frac{1}{n} \sum_{i=1}^n (x_i - \overline{\mu})^2$$

Why?

Density estimation

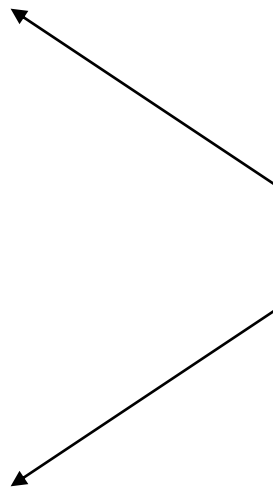
- Binary and discrete variables:

Easy: Just count!

- Continuous variables:

Harder (but just a bit): Fit a model

But what if we only have very few samples?




Important points

- Random variables
- Chain rule
- Bayes rule
- Joint distribution, independence, conditional independence
- MLE


Probability Density Function

- Discrete distributions



$$\sum_i P(X = x_i) = 1$$

- Continuous: Cumulative Density Function (CDF): $F(a)$



$$P(x \leq a) = \int_{-\infty}^a f(\tau) d\tau$$

Cumulative Density Functions

- Total probability $P(\Omega) = \int_{-\infty}^{\infty} f(x)dx = 1$
- Probability Density Function (PDF) $\frac{d}{dx}F(x) = f(x)$


- Properties:

$$P(a \leq x \leq b) = \int_a^b f(x)dx = F(b) - F(a)$$

$$\lim_{x \rightarrow -\infty} F(x) = 0$$

$$\lim_{x \rightarrow \infty} F(x) = 1$$

$$F(a) \geq F(b) \quad \forall a \geq b$$



Expectations

- Mean/Expected Value:

$$E[x] = \bar{x} = \int x f(x) dx$$

- Variance:

$$Var(x) = E[(x - \bar{x})^2] = E[x^2] - (\bar{x})^2$$

- In general:

$$E[x^2] = \int x^2 f(x) dx$$

$$E[g(x)] = \int g(x) f(x) dx$$

Multivariate

- Joint for (x,y)

$$P((x, y) \in A) = \int \int_A f(x, y) dx dy$$

- Marginal:

$$f(x) = \int f(x, y) dy$$

- Conditionals:

$$f(x|y) = \frac{f(x, y)}{f(y)}$$

- Chain rule:

$$f(x, y) = f(x|y)f(y) = f(y|x)f(x)$$

Bayes Rule

- Standard form:

$$f(x|y) = \frac{f(y|x)f(x)}{f(y)}$$

- Replacing the bottom:

$$f(x|y) = \frac{f(y|x)f(x)}{\int f(y|x)f(x)dx}$$

Binomial

- Distribution:

$$x \sim \textit{Binomial}(p, n)$$

$$P(x = k) = \binom{n}{k} p^k (1 - p)^{n-k}$$

- Mean/Var:

$$E[x] = np$$

$$\textit{Var}(x) = np(1 - p)$$

Uniform

- Anything is equally likely in the region $[a,b]$
- Distribution:


$$x \sim U(a, b)$$

- Mean/Var

$$f(x) = \begin{cases} \frac{1}{b-a} & a \leq x \leq b \\ 0 & \text{otherwise} \end{cases}$$

$$E[x] = \frac{a+b}{2}$$

$$Var(x) = \frac{a^2 + ab + b^2}{3}$$



Gaussian (Normal)

- If I look at the height of women in country xx, it will look approximately Gaussian
- Small random noise errors, look Gaussian/Normal

- Distribution:


$$x \sim N(\mu, \sigma^2)$$

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

- Mean/var

$$E[x] = \mu$$

$$Var(x) = \sigma^2$$



Why Do People Use Gaussians

- Central Limit Theorem: (loosely)
 - Sum of a large number of IID random variables is approximately Gaussian

Multivariate Gaussians

- Distribution for vector x

$$x = (x_1, \dots, x_N)^T, \quad x \sim N(\mu, \Sigma)$$

- PDF:

$$f(x) = \frac{1}{(2\pi)^{\frac{N}{2}} |\Sigma|^{\frac{1}{2}}} e^{-\frac{1}{2} (x-\mu)^T \Sigma^{-1} (x-\mu)}$$

$$E[x] = \mu = (E[x_1], \dots, E[x_N])^T$$

$$Var(x) \rightarrow \Sigma = \begin{pmatrix} Var(x_1) & Cov(x_1, x_2) & \dots & Cov(x_1, x_N) \\ Cov(x_2, x_1) & Var(x_2) & \dots & Cov(x_2, x_N) \\ \vdots & & \ddots & \vdots \\ Cov(x_N, x_1) & Cov(x_N, x_2) & \dots & Var(x_N) \end{pmatrix}$$

Multivariate Gaussians

$$f(x) = \frac{1}{(2\pi)^{\frac{N}{2}} |\Sigma|^{\frac{1}{2}}} e^{-\frac{1}{2} (x-\mu)^T \Sigma^{-1} (x-\mu)}$$


$$E[x] = \mu = (E[x_1], \dots, E[x_N])^T$$

$$Var(x) \rightarrow \Sigma = \begin{pmatrix} Var(x_1) & Cov(x_1, x_2) & \dots & Cov(x_1, x_N) \\ Cov(x_2, x_1) & Var(x_2) & \dots & Cov(x_2, x_N) \\ \vdots & & \ddots & \vdots \\ Cov(x_N, x_1) & Cov(x_N, x_2) & \dots & Var(x_N) \end{pmatrix}$$

$$cov(x_1, x_2) = \frac{1}{n} \sum_{i=1}^n (x_{1,i} - \mu_1)(x_{2,i} - \mu_2)$$


Covariance examples

Anti-correlated




Covariance: -9.2

Correlated



Covariance: 18.33

Independent (almost)



Covariance: 0.6

Sum of Gaussians

- The sum of two Gaussians is a Gaussian:

$$x \sim N(\mu, \sigma^2) \quad y \sim N(\mu_y, \sigma_y^2)$$

$$ax + b \sim N(a\mu + b, (a\sigma)^2)$$

$$x + y \sim N(\mu + \mu_y, \sigma^2 + \sigma_y^2)$$